

ECE 448 cheat-sheet

Midterm

sp 2023 2JUI

Expert system = knowledge base
"Fuzzy logic" + logical rules

Failure: The law of out-of-vocabulary word
fragility

rational: maximizing expected utility

pros: Generality; Practicality; Solvability

AI: think and act humanly & rationally

$E \leftarrow$ world states

agent: PEAS $A \leftarrow$ transition model
Actuators

Types of agents: RLGU Utility-Directed

Reflex; Internal-state; Goal-Directed;

Properties of Environments: ODSEC Continuous
Observable; Deterministic; Episodic; Static;

$P \leftarrow$ utility $U_t = f(E_t, A_t; (t-1))$

- Fully observable vs. partially observable
- Deterministic vs. stochastic
- Episodic vs. sequential
- Static vs. dynamic
- Discrete vs. continuous
- Single agent vs. multi-agent
- Known vs. unknown

Spam Filter:

$$E_t = \{M_1, \dots, M_t, L_1, \dots, L_t\}$$

$L_t=1$ if the email is spam, $A_t=1$ if reject it.

$$U_t = -\sum_{i=1}^t C_{FR} A_i (1-L_i) + C_{FA} L_i (1-A_i)$$

Performance measure: An objective criterion
for success of an agent's behavior

RLGU: I(past) G(past+future)
U(past+future + value) sensors \rightarrow AI
↓
effectors

strategic: the environment is deterministic
except for the actions of other agents

Search

State space: the initial state, actions
and transition model

unvisited set:
(directed graph) Frontier

to handle repeated states: use a explored
set. $O\{|S|\}$ ($O\{|S|\} < O\{b^m\}$)

DFS: Frontier is a LIFO stack

T_n and S_n are measured by

b : maximum branching factor of the
search tree

d : depth of the optimal solution

m : maximum length of any path in the
state space (possibly infinite) $O\{b^m\}$

$|S|$: number of distinct states

$O(bm)$

BFS: frontier is FIFO queue
complete and optimal

Time: $O\{b^d\}$ a b-ary tree of depth

Space: $O\{b^d\} \gg O(bm)$

Dijkstra's algorithm: Uniform-cost search frontier is a priority queue ordered by the path cost

Time: $O(b^d)$ Priority queue is $O(\log_2 d)/\text{node}$

Space: $O(b^d)$

Heuristics: accurate if $h(n) \leq d(n)$
cheap if $\ll O(b^d)$

Manhattan distance: $h(n) = |x_n - x_a| + |y_n - y_a|$

Greedy: $\min h(n)$

A*: $f(n) = g(n) + h(n)$ $g(n)$: path cost
 $c(n)$

$h(n)$ is admissible if $d(n) \geq h(n)$

Bad interaction between A* and es.

To solve it, ① Nodes on the explored set are tagged by $c(n)$; ② use a heuristics not only admissible but also consistent

$h(n)$ is consistent if $d(n) - d(p) \geq h(n) - h(p)$

Dijkstra = A* with $h(n) = 0$

Relaxed problem: fewer restrictions

Admissible heuristics $h_1(n)$ and $h_2(n)$

if $h_2(n) \geq h_1(n)$ for all n , then

h_2 dominates h_1

CSP: D possible values with N variables
commutative $N! D^N$

Recursive-backtracking search

CDFS with heuristics that

① choose the next variable to assign

Least Remaining Values (LRV)

Most Constraining Variable (MCV)

② choose a value for that var

Least Constraining Assignment (LCA)

③ Early detection of failure

Forward Checking: Terminate if any illegal

Arc Consistency: Constraint propagation check and re-check

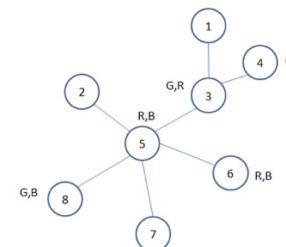
```
function RECURSIVE-BACKTRACKING(assignment, csp)
    if assignment is complete then return assignment
    var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp)
        if value is consistent with assignment given CONSTRAINTS[csp]
            add {var = value} to assignment
            result ← RECURSIVE-BACKTRACKING(assignment, csp)
            if result ≠ failure then return result
            remove {var = value} from assignment
    return failure
```

Apply inference to reduce the space of possible assignments and detect failure early

```
function AC-3(csp) returns the CSP, possibly with reduced domains
inputs: csp, a binary CSP with variables { $X_1, X_2, \dots, X_n$ }
local variables: queue, a queue of arcs, initially all the arcs in csp
```

```
while queue is not empty
    ( $X_i, X_j$ ) ← REMOVE-FIRST(queue)
    if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
        for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
            add ( $X_k, X_i$ ) to queue
```

```
function REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds
removed ← false
for each  $x$  in DOMAIN[ $X_i$ ]
    if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
        then delete  $x$  from DOMAIN[ $X_i$ ]; removed ← true
return removed
```



Solution:

- Sort the nodes into a search order with no crossed edges, e.g., 2,8,7,6,5,3,1,4.
- Perform arc consistency on each pair of nodes, from right to left. With the ordering specified above, the following changes are made: 3 cannot be G. 5 cannot be R. 2, 8, 7, 6 cannot be B.
- Assign the variables to any remaining value, from left to right. For example, 2=G, 8=G, 7=G, 6=R, 5=B, 3=R, 1=B, 4=G.

Algorithm	Complete?	Optimal?	Time complexity	Space complexity	Implement the Frontier as a...
BFS DFS if many sol	Yes No	If all step costs are equal No	$O(b^d)$ $O(b^m)$	$O(b^d)$ $O(bm)$	Queue Stack
IDS	Yes	If all step costs are equal	$O(b^d)$	$O(bd)$	Stack
UCS	Yes	Yes	Number of nodes w/ $g(n) \leq C^*$	Number of nodes w/ $g(n) \leq C^*$	Priority Queue sorted by $g(n)$
Greedy	No	No	Worst case: $O(b^m)$ Best case: $O(bd)$	Worse case: $O(b^m)$ Best case: $O(bd)$	Priority Queue sorted by $h(n)$
A*	Yes	Yes	Number of nodes w/ $g(n) + h(n) \leq C^*$	Number of nodes w/ $g(n) + h(n) \leq C^*$	Priority Queue sorted by $h(n) + g(n)$

Universal Instantiation: $\forall x, F(x)$

Existential instantiation: $\exists x, F(x)$

Forward Chaining: generates a new sentences by combining many diff states.

Backward chaining: apply {set of known sentences} to generate {set of desired sentences} in order to the target sentence

Diff between Planning and Theorem Proving:

Planning may negate some of its preconditions

1st h(n): number of goal sentences left to achieve;

2nd h(n): planning graph

upper bound Ant Snt + Mutex links
until no mutex
(Convergent) h(n) can be stages that
no longer mutex links

Planning is PSPACE-complete

QBF: $\exists x_1 \forall x_2 \exists x_3 \forall x_4 (\wedge_1 V \neg x_3 V x_4) \wedge$

quantified boolean formulae ($\neg x_2 V x_3 V \neg x_4$)
(games)

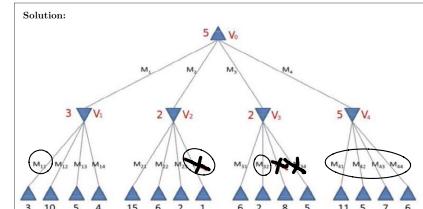
SAT: $\exists x_1 \exists x_2 \exists x_3 \exists x_4 (x_1 V \neg x_2 V x_4) \wedge$
($\neg x_2 V x_3 V \neg x_4$)

puzzles

Two player's game: tic-tac-toe

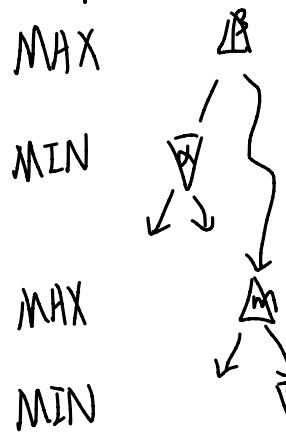
Max-min
me my opponent

zero-sum game
 \Rightarrow minimax tree



Alpha-beta pruning ($O(b^{\frac{m}{2}})$ from $O(b^m)$)
 β is the value of the best choice for the MIN player found so far (lowest number)

α is the value of the best choice for the MAX player found at any choice above node n .



we want to compute

① MIN value at n , loop over n 's children, it decreases

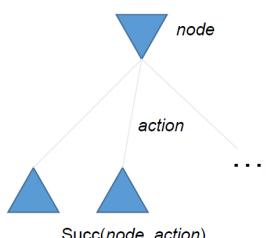
② MAX value at m , loop over m 's children, it increases

$$\alpha \leq \beta$$

Function $action = \text{Alpha-Beta-Search}(node)$
 $v = \text{Min-Value}(node, -\infty, \infty)$
return the action from node with value v

α : best alternative available to the Max player
 β : best alternative available to the Min player

Function $v = \text{Min-Value}(node, \alpha, \beta)$
if Terminal(node) return Utility(node)
 $v = +\infty$
for each action from node
 $v = \text{Min}(v, \text{Max-Value}(\text{Succ}(node, action), \alpha, \beta))$
 if $v \leq \alpha$ return v
 $\beta = \text{Min}(\beta, v)$
end for
return v



branch-factor: number of children nodes for each parent

Function $action = \text{Alpha-Beta-Search}(node)$
 $v = \text{Max-Value}(node, -\infty, \infty)$
return the action from node with value v

α : best alternative available to the Max player
 β : best alternative available to the Min player

Function $v = \text{Max-Value}(node, \alpha, \beta)$
if Terminal(node) return Utility(node)
 $v = -\infty$
for each action from node
 $v = \text{Max}(v, \text{Min-Value}(\text{Succ}(node, action), \alpha, \beta))$
 if $v \geq \beta$ return v
 $\alpha = \text{Max}(\alpha, v)$
end for
return v

③ Transposition table

④ Forward Pruning

⑤ Lookup table

Limited-horizon search

$$\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

Horizon effect: the opponent's move effect

delayed but no avoided

Remedies: ① Quiescence search; ② singular extension

Nash equilibrium: A pair of strategies such that no player can get a bigger payoff by switching strategies, provided that the other player sticks with the same strategy.

Dominant strategy

Pareto optimal outcome

For better: ① Super rationality: Categorical Imperative

② Repeated Games:

Cooperation

↳ in stag Hunt

Sealed-bid second-price auction:
the winner pays the price of the second-highest bid (bid your V)
truth revealing mechanism

Probability: $P(A|B) = \frac{P(A, B)}{P(B)}$ Conditional distribution

$$P(\neg A, B) = P(B) - P(A, B)$$

A, B is independent iff $P(A, B) = P(A) \cdot P(B)$
computational complexity limitation
environment: unknown / partially observable
stochastic

expected utility:

$$E[\text{Utility} | \text{Action}] = \sum_{\text{outcomes}} P(\text{outcome} | \text{action}) \cdot \text{Utility}(\text{outcome})$$

Solution: The stag hunt is a coordination game. The payoff matrix is

	Player 1: Cooperate	Player 1: Defect
Player 1: Cooperate	2, 2	0, 1
Player 2: Defect	1, 0	1, 1

There are no dominant strategies. The pure-strategy Nash equilibria are (Cooperate, Cooperate) and (Defect, Defect). The Pareto-optimal solution is (Cooperate, Cooperate).

The game of Chicken is an anti-coordination game. The payoff matrix is

	Player 1: Chicken	Player 1: Drive
Player 1: Chicken	0, 0	-1, 1
Player 2: Drive	1, -1	-10, -10

There are no dominant strategies. The Pareto optimal solutions are (Chicken, Chicken), (Chicken, Drive), and (Drive, Chicken). The pure-strategy Nash equilibria are (Chicken, Drive) and (Drive, Chicken).

Solution: Minimax solution maximizes, over all of your possible actions, the minimum, over all of your opponent's possible actions, of your reward.

- Prisoner's Dilemma: Defect. Result is also the Nash equilibrium.
- Stag Hunt: take the Hare. Result is one of the two Nash equilibria.
- Game of Chicken: chicken out. Result is not a Nash equilibrium.

Nash equilibrium assumes that you know what the other player will do, and can respond appropriately. Minimax makes more sense if you want to limit your losses, and have no way to predict the other player's behavior.

utility theory is used to represent and infer preference
decision theory = probability theory + utility theory

Frequentism, Subjectivism

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$

Atomic events: mutually exclusive and exhaustive

$$\text{joint distribution: } \sum P_i = 1$$

marginal distribution: from $P(X, Y)$ to get

$$P(X) \text{ and } P(Y)$$

$$P(A, B) = P(A|B) \cdot P(B) \quad \text{product rule}$$

$$= P(B|A) \cdot P(A) \quad \text{chain rule}$$

$$P(A, B, \dots, N)$$

$$\Rightarrow P(A|B) = P(A) \quad \text{mutually exclusive}$$

$$P(B|A) = P(B) \quad P(A \vee B) = P(A) + P(B)$$

conditionally independent

$$P(A \wedge B | C) = P(A|C) \cdot P(B|C)$$

Expect Minimax

$$\max(\min(E(\text{Reward})))$$

SL & RL

Monte Carlo Tree Search